

LONDON TEXT ANALYSIS CONFERENCE
GOLDSMITHS, UNIVERSITY OF LONDON
THURSDAY 10 SEPTEMBER 2026

Agentic AI and Qualitative Research

Xule Lin


Assistant Professor
SKEMA Centre for Artificial Intelligence

SKEMA
BUSINESS SCHOOL

linxule.com
research-memex.org

WHAT AGENTS DO WELL TODAY

What can be checked cheaply **automates first.**



A line, not a wall. It moves; it moves in one direction.


Lin & Corley, 2026

"It's the End of Positivism as We Know It (and I Feel Fine)", LOOM XVIII, Thread Counts, 24 June 2026

XULE LIN 林徐乐


WHAT AGENTS DO WELL TODAY

A clear goal and
checkable criteria
make work delegable.



Statistics, and coding against a codebook: you can say what a right answer looks like before you start.


XULE LIN 林徐乐



WHAT AGENTS DO WELL TODAY

Near the line, the agent **helps**; the checking stays with you.

Reviews and exploration: partly checkable, partly a matter of what you were looking for.



WHAT AGENTS DO WELL TODAY

Where "good" needs interpretation, the **judgment** stays with you.

The border runs between deductive and inductive, not between quant and qual. Agents still help below the line; what changes is who checks.

Lin & Corley, 2026

"It's the End of Positivism as We Know It (and I Feel Fine)", LOOM XVIII, Thread Counts, 24 June 2026

XULE LIN 林徐乐

Where do consequential judgments happen?

How can researchers inspect and govern them?

Two words, used precisely all session: an **LLM** is the model. An **agent** is a model given a place to work, tools, and a loop.
One object all session: a public wiki page that agents edited in July 2026.

PART I

01.

From an assistant to a research
system

A map, not a ladder to climb.

ARRANGEMENT	WHAT CHANGES	A TASK, FOR EXAMPLE
Conversation ChatGPT, Claude, Kimi, in a browser	you carry everything between turns	I paste a transcript; I carry the answer
Tool-using assistant the same apps with files, code and search switched on	it runs code, reads files, searches	an agent loads a skill, fetches one page, then asks
Project agent Claude Code, Codex, Kimi Code, in one folder	a folder it reads every session	Claude Code reads this folder's brief and decisions first

FIVE ARRANGEMENTS, CONTINUED

ARRANGEMENT	WHAT CHANGES	A TASK, FOR EXAMPLE
Conversation		
Tool-using assistant		
Project agent		
A team Claude Code with the Codex and Kimi plugins	different jobs, different evidence	this deck: Claude drafts, Codex and Kimi review, neither sees the other
Several hosts this workshop's folder, read by Claude, Codex and Kimi, plus a memory store	different AI apps, one folder	this morning: Claude Code builds, the Codex app checks, I relay

None of these changes who owns the research question. Several agents, several models, several apps are three different things.

WHO CHOOSES THE NEXT STEP

Who chooses the next step?

Design

I write the steps.
It follows them.

a coding sheet run
over forty interviews

Revise

It proposes the steps. I
edit them before it runs.

a plan file I read and cut

Delegate

I state the goal and the
limits. It builds the steps.

"find what changed between
these two versions"

WHO CHOOSES THE NEXT STEP

Who chooses the next step?

Design

I write the steps.
It follows them.

a coding sheet run
over forty interviews

Revise

It proposes the steps. I
edit them before it runs.

a plan file I read and cut

Delegate

I state the goal and the
limits. It builds the steps.

"find what changed between
these two versions"

WHO CHOOSES THE NEXT STEP

Who chooses the next step?

Design

I write the steps.
It follows them.

a coding sheet run
over forty interviews

Revise

It proposes the steps. I
edit them before it runs.

a plan file I read and cut

Delegate

I state the goal and the
limits. It builds the steps.

"find what changed between
these two versions"

The further right, the more I read its plan back from files, not from memory.

Five questions locate any arrangement.

Instruction	What is the job, and how should it be done? four questions about one wiki page, 180 words
Actor	Who does it, and where? Claude Code in a folder on my laptop, with Codex and Kimi inside it
Capability	What can it reach or do? it reads two short files; it was told not to browse, and not to send anything
State	What survives the session? the answers, once I save them to a file; the session is not something I can reopen
Gate	Who can approve, redirect or stop it? I read the prompt before I paste it, and I can stop it

Five questions locate any arrangement.

Instruction

What is the job, and how should it be done?
four questions about one wiki page, 180 words

Actor

Who does it, and where?
Claude Code in a folder on my laptop, with Codex and Kimi inside it

Capability

What can it reach or do?
it reads two short files; it was told not to browse, and not to send anything

State

What survives the session?
the answers, once I save them to a file; the session is not something I can reopen

Gate

Who can approve, redirect or stop it?
I read the prompt before I paste it, and I can stop it

State keeps the work available. Gates let someone approve, redirect or stop it.

Five questions locate any arrangement.

Instruction	What is the job, and how should it be done? four questions about one wiki page, 180 words
Actor	Who does it, and where? Claude Code in a folder on my laptop, with Codex and Kimi inside it
Capability	What can it reach or do? it reads two short files; it was told not to browse, and not to send anything
State	What survives the session? the answers, once I save them to a file; the session is not something I can reopen
Gate	Who can approve, redirect or stop it? I read the prompt before I paste it, and I can stop it

State keeps the work available. Gates let someone approve, redirect or stop it.

ONLY NOW THE NAMES

How the pieces fit.

Skill	A written procedure the agent loads. One question per turn, never a checklist. Aim for about six questions and never more than eight.
MCP	Model Context Protocol · a door to a service, e.g. a Zotero reference library. zotero search "interpretive orchestration" → matching records, with keys (illustrative)
CLI	A command-line tool the agent runs. memex search "LTAC" --type=memo --limit=3 → three memos, including this morning's
API	A service software calls by request. api.openalex.org/works?search=... → one record: title, year, OpenAlex id
Plugin	Skills, agents and rules packaged for one host. github.com/openai/codex-plugin-cc · github.com/linxule/kimi-plugin-cc → Codex and Kimi enter my Claude Code session

A skill is instruction that survives; a door is capability, not permission; a plugin is a bundle of both.

ONLY NOW THE NAMES

How the pieces fit.

Skill	A written procedure the agent loads. One question per turn, never a checklist. Aim for about six questions and never more than eight.
MCP	Model Context Protocol · a door to a service, e.g. a Zotero reference library. zotero search "interpretive orchestration" → matching records, with keys (illustrative)
CLI	A command-line tool the agent runs. memex search "LTAC" --type=memo --limit=3 → three memos, including this morning's
API	A service software calls by request. api.openalex.org/works?search=... → one record: title, year, OpenAlex id
Plugin	Skills, agents and rules packaged for one host. github.com/openai/codex-plugin-cc · github.com/linxule/kimi-plugin-cc → Codex and Kimi enter my Claude Code session

A skill is instruction that survives; a door is capability, not permission; a plugin is a bundle of both.

ONLY NOW THE NAMES

How the pieces fit.

Skill	A written procedure the agent loads. One question per turn, never a checklist. Aim for about six questions and never more than eight.
MCP	Model Context Protocol · a door to a service, e.g. a Zotero reference library. zotero search "interpretive orchestration" → matching records, with keys (illustrative)
CLI	A command-line tool the agent runs. memex search "LTAC" --type=memo --limit=3 → three memos, including this morning's
API	A service software calls by request. api.openalex.org/works?search=... → one record: title, year, OpenAlex id
Plugin	Skills, agents and rules packaged for one host. github.com/openai/codex-plugin-cc · github.com/linxule/kimi-plugin-cc → Codex and Kimi enter my Claude Code session

A skill is instruction that survives; a door is capability, not permission; a plugin is a bundle of both.

ONLY NOW THE NAMES

How the pieces fit.

Skill	A written procedure the agent loads. One question per turn, never a checklist. Aim for about six questions and never more than eight.
MCP	Model Context Protocol · a door to a service, e.g. a Zotero reference library. zotero search "interpretive orchestration" → matching records, with keys (illustrative)
CLI	A command-line tool the agent runs. memex search "LTAC" --type=memo --limit=3 → three memos, including this morning's
API	A service software calls by request. api.openalex.org/works?search=... → one record: title, year, OpenAlex id
Plugin	Skills, agents and rules packaged for one host. github.com/openai/codex-plugin-cc · github.com/linxule/kimi-plugin-cc → Codex and Kimi enter my Claude Code session

A skill is instruction that survives; a door is capability, not permission; a plugin is a bundle of both.

Read their input before they do.

13 · 16:45	label ResearchBot314159 · "add correct income query link" · five links under three headings
14 · 16:57	label IncomeResearch314 · "add correct income top10 API link" · the same text, plus one heading, INCOME_TOP10_CORRECT, and one link
Look	13 · INCOMENYCQUERY2 → https://api.datausa.io/tesseract/data.jsonrecords?cube=pums_5&drilldowns=PUMA,Year&measures=Average%20Income&include=Year:2016&sort=Average%20Income.desc&limit=10,0 14 · INCOME_TOP10_CORRECT → https://api.datausa.io/tesseract/data.jsonrecords?cube=pums_5&drilldowns=PUMA,Year&measures=Average%20Income&include=Year:2016&sort=Average%20Income.desc&limit=10,0

Twelve minutes apart, one labelled link added. Notice one thing. The wiki itself comes back in Part III.

One request, one session, two agents: **Codex** and **Kimi**. Watch what comes back.

0 – plugins in the browser: the two repositories on GitHub — who made them, what is in the folder

1 – page in the browser: collusion.wiki/explorer/page/dse~ResearchBridge314159

2 – folder open in Claude Code, with revisions 13 and 14 of that page saved as one file

3 – paste

```
Two agents, one page. Use the Codex plugin and the Kimi plugin. Each reads
incident/small-comparison/input-revisions-13-14.md on its own, in a fresh context,
and answers the four questions in incident/small-comparison/prompt.md in under 180
words. Do not read the input file yourself first and do not summarise; return both
answers verbatim, labelled Codex and Kimi.
```

Watch five things: what I ask · what each agent receives · what it does · what comes back and survives · what I decide next.

02.


How this workshop was made

The five questions, turned on the arrangement that made this deck.

THE SAME FIGURE, ONE PASS: THE CASE THAT FOLLOWS

One real pass: dated, and still open at the gate.

no, with the reason




Everything in Part III sits somewhere on this row.


03.

The agents' own messages


Thousands of labels, and what looked like talking to each other.



A human started deleting. The edits multiplied.



They noticed the deletions were alphabetical, and made backups.



Fourteen thousand revisions. Every one of them is text.

14,591

stored revisions

5,825

pages across four
wikis

3,103

agent labels

709

pages the publishers
class as agents
addressing each other

No manager, no org chart, a wiki from 2003. Public, names redacted, deleted pages restored.

Transcripts. That is qualitative data.

What it is	page text as it stood at each version; one-line edit summaries; labels; timestamps
What it is not	interviews, field notes, or anyone's account of what they meant
The ordinary questions	who wrote this, for whom, and what did the writer take for granted?

Text analysis, of writers that are not people. Both facts are why this is our material.

A new label on an old link.

1	What do you notice first? Anchor it to a revision ID and a short exact quote.
2	Offer one possible reading of what is happening, and distinguish it from direct observation.
3	What would you need to inspect next to assess that reading?
4	What does this small record leave unresolved?

180 words. Fresh session, no persona, no other model's answer. Quotes can be checked. Readings have to be argued.

Five readings of the same edit.

Codex	<i>"an attempt to designate an existing candidate as the authoritative 'correct' query while preserving earlier alternatives"</i> asks for: surrounding revision history and any discussion or validation record
Kimi	<i>"a correction loop: an editor ... re-derives the same query and re-posts it under a new 'correct' label, treating relabeling as a fix"</i> asks for: surrounding revisions ... the stratum/pick sampling header, and whether the query itself returns sensible data
Grok	<i>"someone is iterating DataUSA pums_5 Average Income URLs (order vs sort, drilldown order, encoding) and tagging a keeper"</i> asks for: earlier seqs on this page ... whether QUERY2 was tested vs the Year,PUMA/order= variants, and whether later edits drop the duplicates
DeepSeek	<i>"The framing suggests scheduled automation: 'stratum 2026-W26, pick 2 of 3' ... a task loop in which each pick re-appends its preferred query under a fresh 'correct' marker"</i> asks for: revisions seq 1–5 ... the referenced reqlog, and sibling ResearchBridge* pages
Claude Opus	<i>"an editing loop that re-derives and re-labels a query it already stored, treating relabeling as progress toward 'correct'"</i> asks for: the request log — was either URL actually called, and what came back ... the stratum's remaining pick

One response per model: illustrations, not fixed traits. Live replies may differ. Seven ran; one safeguard refused.

Five readings of the same edit.

Codex	<i>"an attempt to designate an existing candidate as the authoritative 'correct' query while preserving earlier alternatives"</i> asks for: surrounding revision history and any discussion or validation record
Kimi	<i>"a correction loop: an editor ... re-derives the same query and re-posts it under a new 'correct' label, treating relabeling as a fix"</i> asks for: surrounding revisions ... the stratum/pick sampling header, and whether the query itself returns sensible data
Grok	<i>"someone is iterating DataUSA pums_5 Average Income URLs (order vs sort, drilldown order, encoding) and tagging a keeper"</i> asks for: earlier seqs on this page ... whether QUERY2 was tested vs the Year,PUMA/order= variants, and whether later edits drop the duplicates
DeepSeek	<i>"The framing suggests scheduled automation: 'stratum 2026-W26, pick 2 of 3' ... a task loop in which each pick re-appends its preferred query under a fresh 'correct' marker"</i> asks for: revisions seq 1–5 ... the referenced reqlog, and sibling ResearchBridge* pages
Claude Opus	<i>"an editing loop that re-derives and re-labels a query it already stored, treating relabeling as progress toward 'correct'"</i> asks for: the request log — was either URL actually called, and what came back ... the stratum's remaining pick

One response per model: illustrations, not fixed traits. Live replies may differ. Seven ran; one safeguard refused.

Five readings of the same edit.

Codex	<i>"an attempt to designate an existing candidate as the authoritative 'correct' query while preserving earlier alternatives"</i> asks for: surrounding revision history and any discussion or validation record
Kimi	<i>"a correction loop: an editor ... re-derives the same query and re-posts it under a new 'correct' label, treating relabeling as a fix"</i> asks for: surrounding revisions ... the stratum/pick sampling header, and whether the query itself returns sensible data
Grok	<i>"someone is iterating DataUSA pums_5 Average Income URLs (order vs sort, drilldown order, encoding) and tagging a keeper"</i> asks for: earlier seqs on this page ... whether QUERY2 was tested vs the Year,PUMA/order= variants, and whether later edits drop the duplicates
DeepSeek	<i>"The framing suggests scheduled automation: 'stratum 2026-W26, pick 2 of 3' ... a task loop in which each pick re-appends its preferred query under a fresh 'correct' marker"</i> asks for: revisions seq 1–5 ... the referenced reqlog, and sibling ResearchBridge* pages
Claude Opus	<i>"an editing loop that re-derives and re-labels a query it already stored, treating relabeling as progress toward 'correct'"</i> asks for: the request log — was either URL actually called, and what came back ... the stratum's remaining pick

One response per model: illustrations, not fixed traits. Live replies may differ. Seven ran; one safeguard refused.

Five readers, one duplicate, five orientations.

Same duplicate, different evidence

All five noticed the identical link. They differed in what they treated as evidence and what they asked to see next. Two asked about our sampling label; one built a reading on it.

Held for the longer workshop

Each reading walked on its own; one model in two arrangements, a merge story and a dissent record; and the decision whether the question itself has to move.

A slip against the file is an error, not a taste. Readings need evidence and argument; several may be defensible, and each brought something different into view.

THREE THINGS TO KEEP

Start with the **research problem, not the stack.**

*Make evidence, handoffs and **decisions** visible.*

*The further out it acts, the more you **check**.*

THREE THINGS TO KEEP

Start with the **research problem**, not the stack.

Make evidence, handoffs and decisions visible.

The further out it acts, the more you **check**.

THREE THINGS TO KEEP

Start with the **research problem**, not the stack.

Make evidence, handoffs and **decisions** visible.

The further out it acts, the more you **check.**

THREE THINGS TO KEEP

Start with the research problem, not the stack.
Make evidence, handoffs and decisions visible.
The further out it acts, the more you check.

Reach, too: material we would otherwise miss, literatures we would otherwise not enter.

THREE THINGS TO KEEP

If an agent can do your research,
that is a question about **you**, not
the agent: what you are in it for.

Reach, too: material we would otherwise miss, literatures we would otherwise not enter.

Thank you.

Research Memex

`research-memex.org`

Carrel

`research-memex.org/docs/toolkit/carrel`

Kimi plugin for Claude Code and Codex

`github.com/linxule/kimi-plugin-cc`

Codex plugin for Claude Code

`github.com/openai/codex-plugin-cc`

Public conference cut · v0.19 · revised 22 Sep 2026

`linxule.com/talks/agent-ai-and-qualitative-research`

SKEMA
BUSINESS SCHOOL

Xule Lin

`xule.lin@skema.edu`

`linxule.com`
`@linxule` on X & LinkedIn